

Live Session Week 1

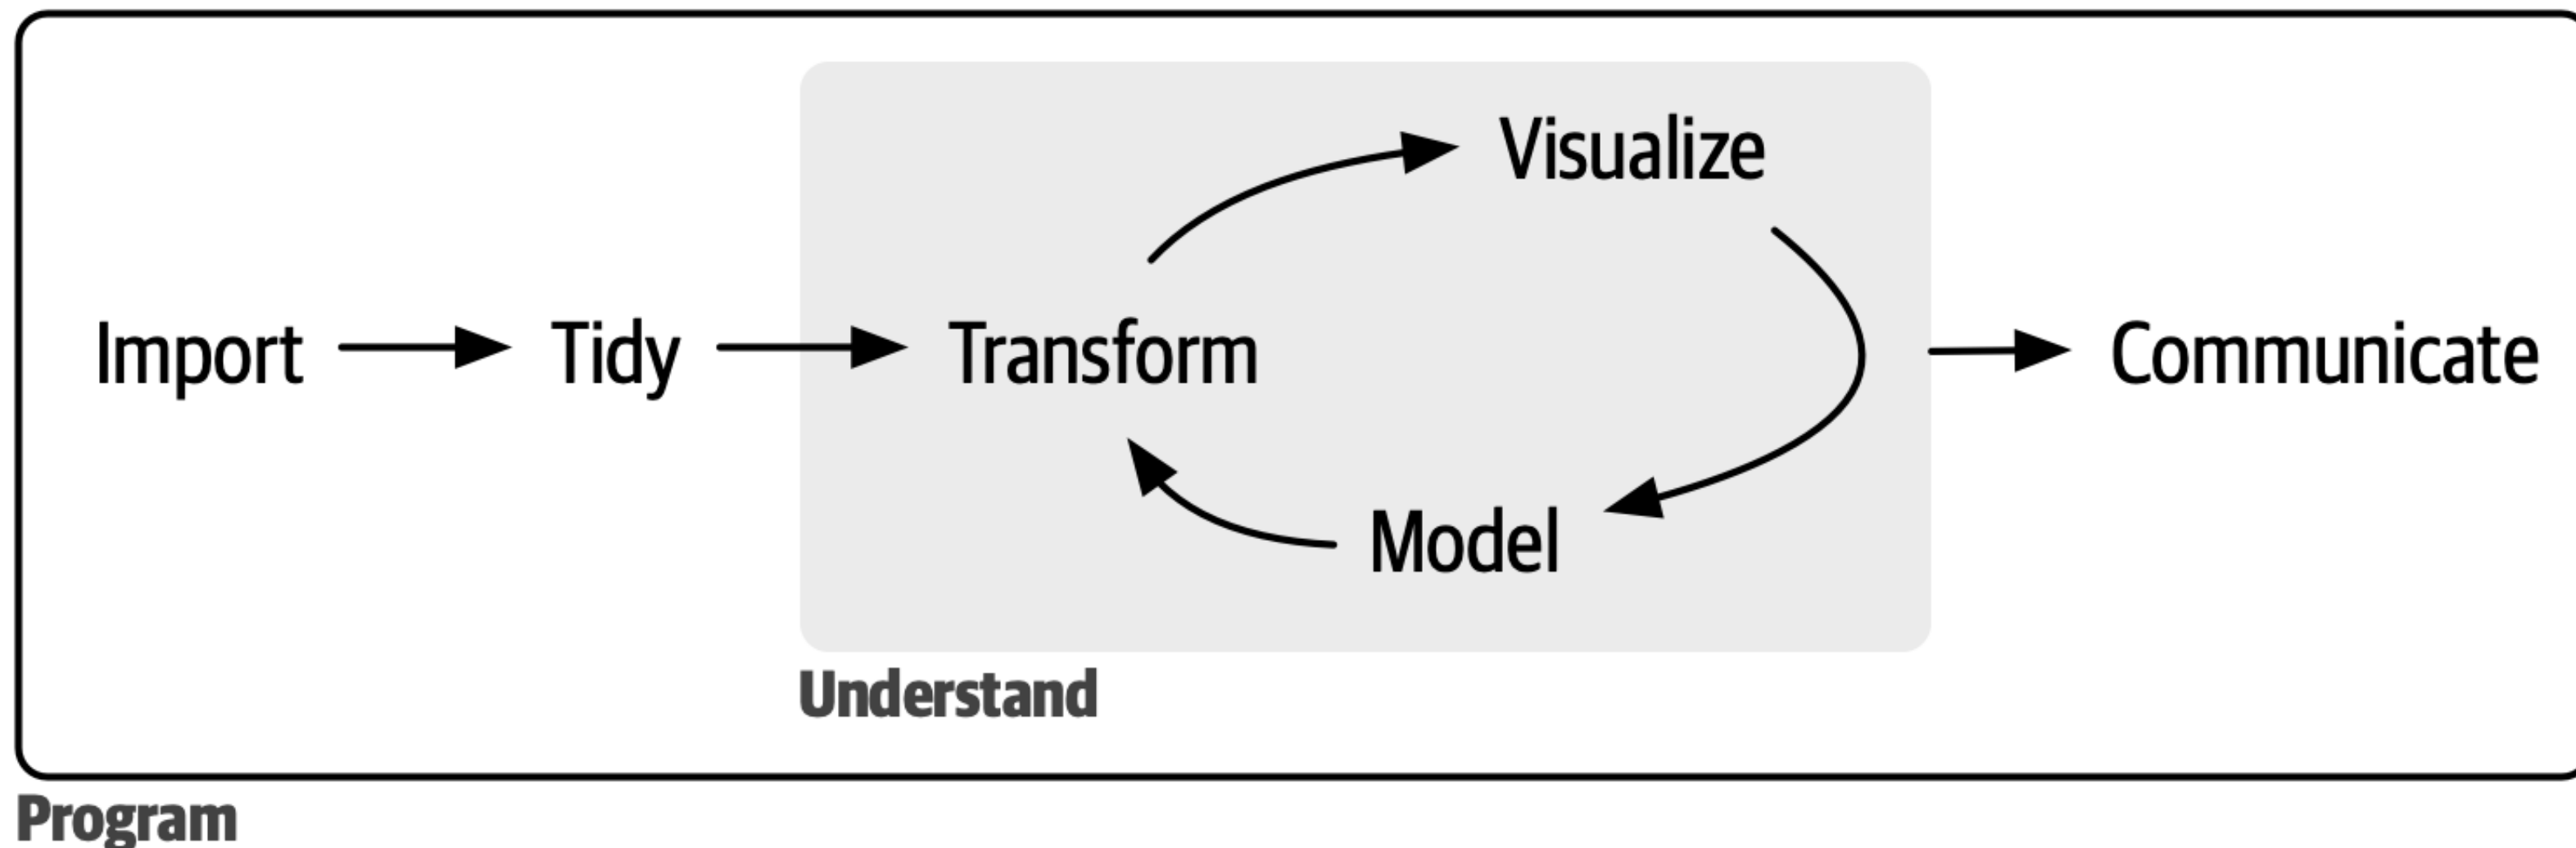
Data Science Workflows

Dr Zak Varty

Outline


1. Review
2. Discussion
3. Break
4. Minimal R Package

Life Cycle of A Data Science Project

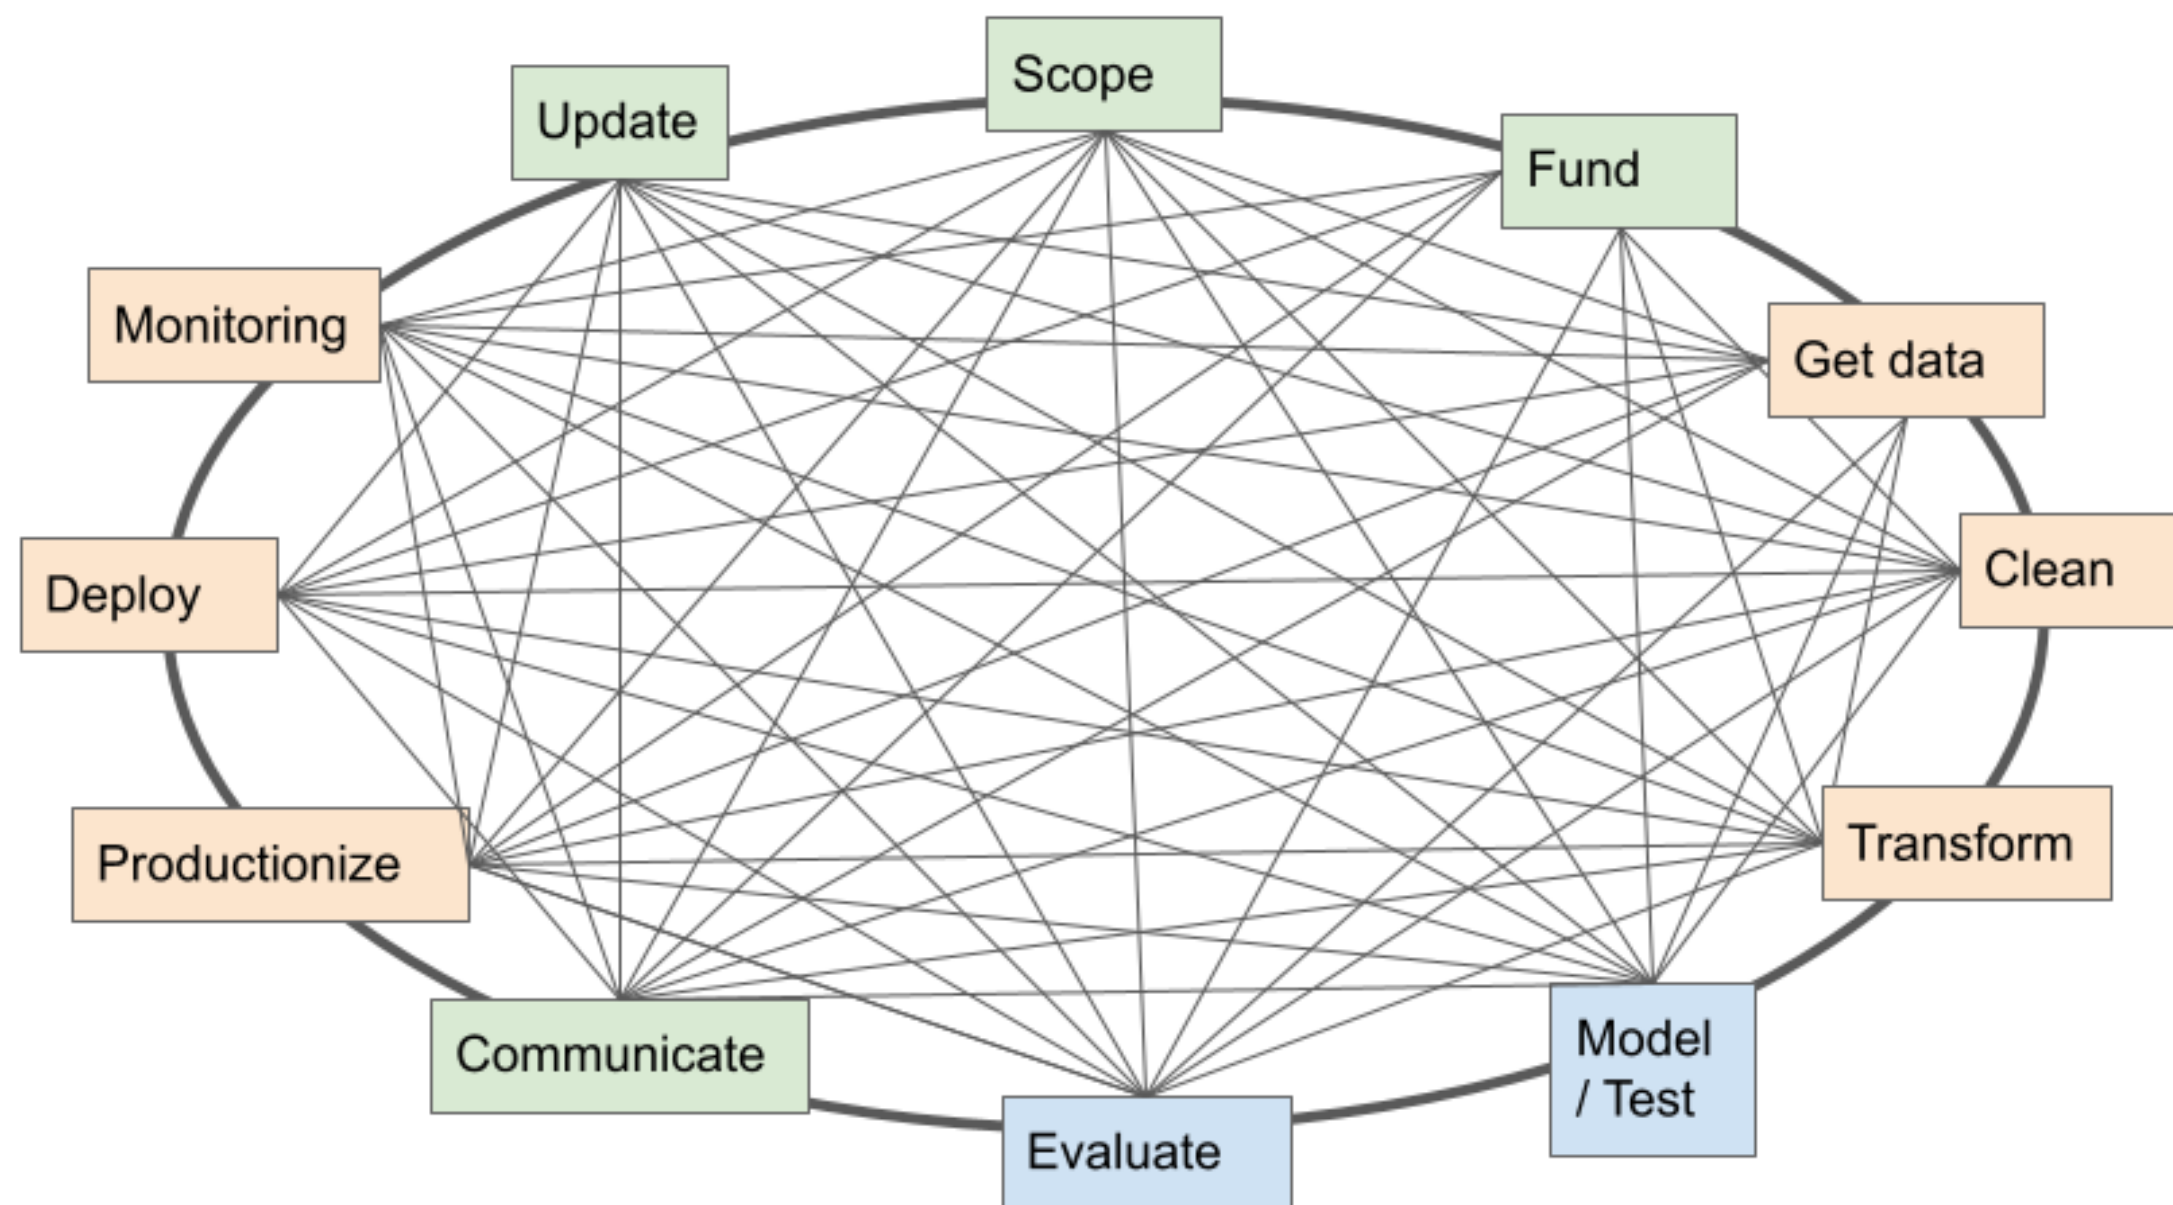


source: R4DS

Life Cycle of A Data Science Project



Life Cycle of A Data Science Project



Task Review

Review - Organising your work

Project management spans all aspects of the data science cycle.

- **Organising your directories:**
 - Benefits of a common directory structure
 - Example project structure
- **Organising your code:**
 - Functional vs Object Oriented Programming
 - Naming conventions and style guide (verb & nouns)
- **Organising your files:**
 - Name with humans, computers and ordering in mind.
 - Pick file types with care.

Review: Task 1 - Github Show and Tell

I asked you to: Find 3 data science projects on Github and explore how they organise their work.

In groups of 2-3, each pitch one of the projects you found.

- What does the project do?
- Who made it?
- Why do you think it is interesting?
- What did you learn by looking at it?

Review: Task 2 - Make Project Templates

I asked you to: Create your own project directory (or directories) for this course and its assignments.

<https://www.menti.com/alrej2iedgtr>



Find a person you haven't spoken to today and explain your reasoning.

Review: Task 3 - Function writing

I asked you to: write a function to calculate the rolling arithmetic mean of a numeric vector.

Pair up with a third person that you haven't spoken to yet today.

Discuss your thought process and compare code.

- What did you have to consider when writing this function?
- Report back: One decision that you made differently **or** a decision that you made the same, but implemented differently.

My thought process (1/n)

```
1 rolling_mean <- function(x){}
```

- Do I pick window length or does the user? (User)
 - What values should I allow? (Integers > 1.)
- Gah, need two inputs: **x** = Vector to smooth, **window_length** = # obs in rolling window. Are these good names?

```
1 rolling_mean <- function(x, window_length){}
```

My thought process (2/n)


- Is the window centred, left-aligned or right aligned?
 - centred for smoothing , right for prediction
 - would I ever want left-aligned?

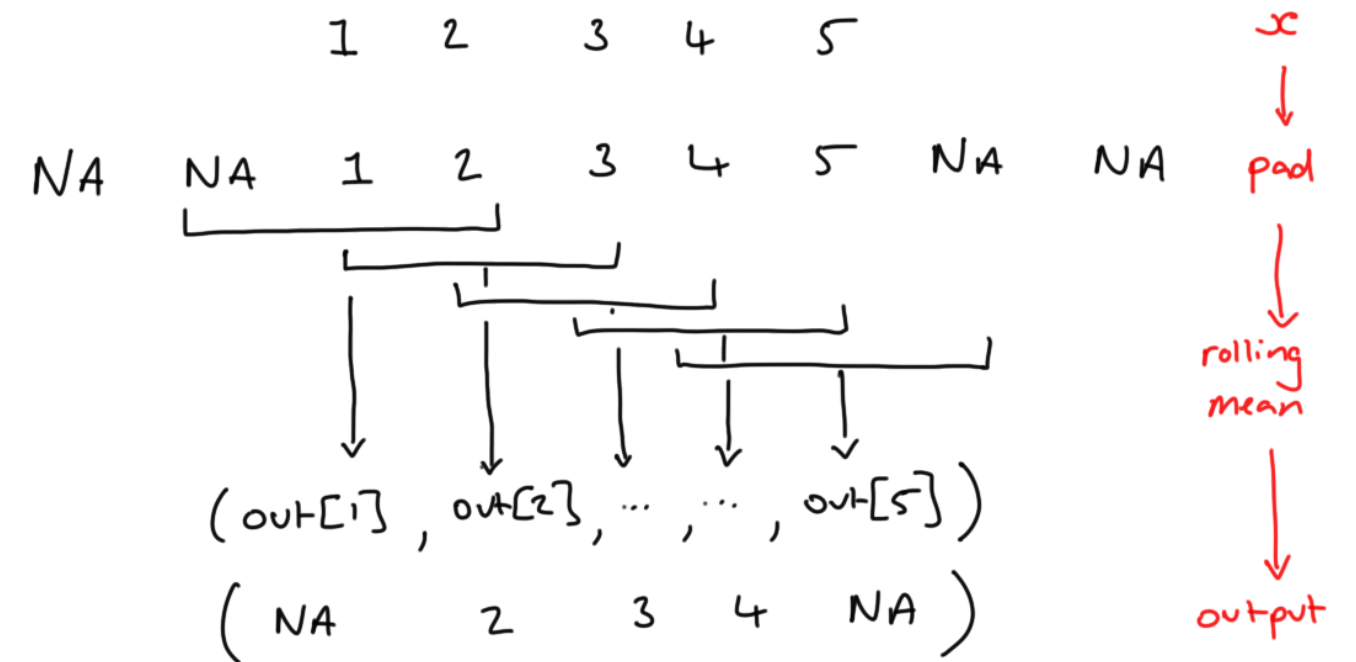
```
1 rolling_mean <- function(x, window_length, window_position = "centre"){}
```

- **window_position** as an argument should this be three separate functions?
 - One function should do one job well. Focus on centred.
 - Can always make a wrapper function later.

```
1 rolling_mean <- function(x, window_length){}
```

My thought process (3/n)

- What do I do at the edges?
 - Miss them and return a shorter vector
 - Average fewer terms
 - Pad with NAs 



Let the worrying commence (4/n)

- What if `window_length` is even?
 - Stupid. Left and right would have worked. Regret.
 - Do I want to return a data frame? Something else?
 - Limit to non-negative odd numbers.
 - How do I do that? `is.integer()`? Nope.

Let the worrying commence (5/n)

- How could this go wrong?
 - Could give a non-vector for **x**
 - Could give a non-numerical vector for **x**
 - Could give a vector of values for **window_width**
 - What happens if they give nothing or **NULL**?
 - What else have I missed ...

The final code

```

1  rolling_mean <- function(x, window_width, ...){
2    # -----Input Checks -----
3    # Check that x is a vector with numerical interpretation
4    stopifnot(is.logical(x) | is.integer(x) | is.double(x) | is.complex(x))
5    stopifnot(length(x) > 0)
6
7    # Check window_width is an odd, positive integer
8    stopifnot(length(window_width) == 1)
9    stopifnot(window_width %% 1 == 0)
10   stopifnot((window_width / 2) %% 1 != 0)
11   stopifnot(window_width > 0)
12
13   # ----- Function Body -----
14
15   # number of values left and right to include in each mean
16   half_width <- floor(window_width / 2)
17   x_padded <- pad_with_NAs(x, n_left = half_width, n_right = half_width)
18   evaluation_locations <- seq_along(x) + half_width
19
20   output <- rep(NA, length(x))

```


My thought process: documenting (5/n)

```
1 #' Calculate the rolling mean of a vector
2 #'
3 #' @param x Vector of values that can be interpreted numerically.
4 #' @param window_width The number of values included in each mean calculation. Should be an odd, positive integer.
5 #' @param ... Additional arguments to pass to the mean() function call.
6 #'
7 #' @return A vector of rolling mean values of the same length as `x`.
8 #' @export
9 #'
10 #' @examples
11 #'
12 #' rolling_mean(x = 1:5, window_width = 3)
13 #' rolling_mean(x = 1:5, window_width = 5)
14 #' rolling_mean(x = 1:5, window_width = 7)
15 #' rolling_mean(x = c(TRUE, TRUE, TRUE, FALSE, TRUE, TRUE, TRUE), window_width = 3)
16 #'
17 rolling_mean <- function(x, window_width, ...){}
```

BREAK

05:00

Minimal R project in 1 hour (or less)

