

# Scalability

Preparing for Production

Dr Zak Varty

# Scalability and Production

When put into production code gets used more and on more data.

We will likely have to consider scalability of our methods in

- Computation time
- Memory requirements

When doing so we have to balance a trade-off between development costs and usage costs.

# Example: Bayesian Inference

- MCMC originally takes ~24 hours
- Identifying and amending bottlenecks in code reduced this to ~24 minutes.


## Is this actually better?

- human hours invested
- frequency of use
- safe / stable / general / readable
- trade for scalability

# Knowing when to worry

Sub-optimal optimisation can be worse than doing nothing

... programmers have spent far too much time worrying about efficiency in **the wrong places** and at **the wrong times**; premature optimisation is the root of all evil (or at least most of it) in programming. - Donald Knuth



# This Lecture

- Basic profiling to find bottlenecks.
- Some simple solutions
- Strategies for scalable (R) code
- Signpost advanced methods & further reading

# Profiling your code: basics

## R as a stopwatch

```
1 t_start <- Sys.time()
2 Sys.sleep(0.5) # YOUR CODE
3 t_end <- Sys.time()
4
5 t_end - t_start
```

Time difference of 0.5098979 secs

```
1 library(tictoc)
2
3 tic()
4 Sys.sleep(0.5) # YOUR CODE
5 toc()
```

0.505 sec elapsed

## With `{tictoc}` we can get fancy

```
1 tic("total")
2 tic("first, easy part")
3 Sys.sleep(0.5)
4 toc(log = TRUE)
5 ## first, easy part: 0.505 sec elapsed
6 tic("second, hard part")
7 Sys.sleep(3)
8 toc(log = TRUE)
9 ## second, hard part: 3.008 sec elapsed
10 toc()
11 ## total: 3.519 sec elapsed
```

# Profiling your code in detail

To diagnose scaling issues you have to understand what your code is doing.


- Stop the code at time  $\tau$  and examine the **call-stack**.
  - The current function being evaluated, the function that called that, the function that called that, ..., top level function.
- Do this a lot and you can measure (estimate) the proportion of working memory (RAM) uses over time and the time spent evaluating each function.

# Profiling: Toy Example

```

1 h <- function() {
2   profvis::pause(1)
3 }
4
5 g <- function() {
6   profvis::pause(1)
7   h()
8 }
9
10 f <- function() {
11   profvis::pause(1)
12   g()
13   profvis::pause(1)
14   h()
15 }

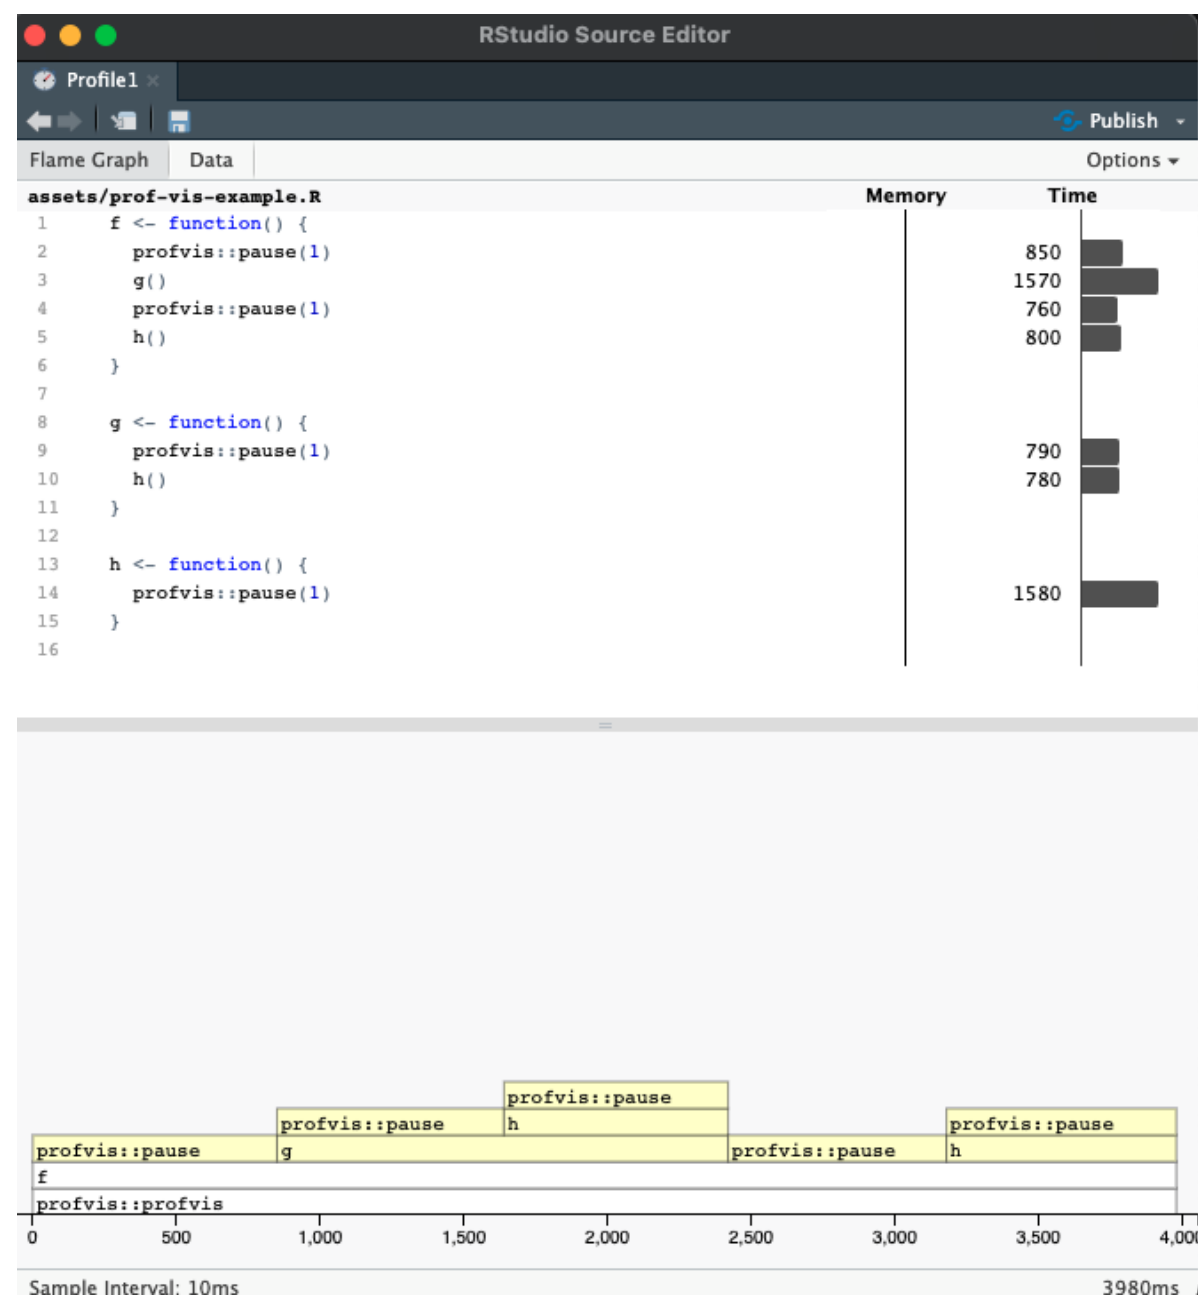
```





# Profiling: How To

```
1 source("assets/prof-vis-example.R")
2 profvis::profvis(f())
```



# Notes on Time Profiling

- Will get slightly different results each time you run the function
  - Changes to internal state of computer
  - Usually not a big deal, mainly effects fastest parts of code
  - Be careful with stochastic simulations
  - Use `set.seed()` to make a fair comparison over many runs.

# Notes on Profiling

## Function Source

```
1 pad_with_NAs  
function(x, n_left, n_right){  
  c(rep(NA, n_left), x, rep(NA, n_right))  
}
```

## Compiled Function

```
1 mean  
function (x, ...)   
UseMethod("mean")  
<bytecode: 0x7fabbac6d718>  
<environment: namespace:base>
```



- Compiled functions have no R source code.
- Profiler does not extend into compiled code, see [{jointprof}](#) if you really need this.

# Memory Profiling

`profvis()` can similarly measure the memory usage of your code.

```
1 x <- integer()
2 for (i in 1:1e4) {
3   x <- c(x, i)
4 }
```

- Copy-on-modify behaviour makes growing objects slow.
- Pre-allocate storage where possible.
- Strategies and structures, see [R inferno](#) and [Efficient R](#).



# Tips to work at scale

# Vectorise

```
1 x <- 1:10
2 y <- 11:20
3 z <- rep(NA, length(x))
4
5 for (i in seq_along(x)) {
6   z[i] <- x[i] * y[i]
7 }
```

```
1 x <- 1:10
2 y <- 11:20
3 z <- x * y
```

Use and write functions with vectorised inputs.

```
1 rnorm(n = 100, mean = 1:10, sd = rep(1, 10))
```

Be careful of recycling!

# Special vectors: Linear Algebra

```
1 x <- diag(x = c(2, 0.5))
2 y <- matrix(data = c(1, 1), ncol = 1)
3
4 x %*% y
```

```
      [,1]
[1,]  2.0
[2,]  0.5
```

More on vectorising: [Noam Ross Blog Post](#)

# For loops in disguise: the apply family

Functional programming equivalent of a for loop. [`apply()`, `mapply()`, `lapply()`, ...]

Apply a function to each element of a list-like object.

```
1 A <- matrix(data = 1:12, nrow = 3, ncol = 4)
2 A
```

```
      [,1] [,2] [,3] [,4]
[1,]    1    4    7   10
[2,]    2    5    8   11
[3,]    3    6    9   12
```

```
1 # MARGIN = 1 => rows, MARGIN = 2 => columns
2 apply(X = A, MARGIN = 1, FUN = sum)
```

```
[1] 22 26 30
```

Generalises functions from `{matrixStats}`

```
1 rowSums(A)
```

```
[1] 22 26 30
```



# For loops in disguise: purrr::map

Iterate over a single object with **map()**:

```
1 mu <- c(-10, 0, 10)
2 purrr::map(.x = mu, .f = rnorm, n = 5)
```

```
[[1]]
[1] -10.630941 -10.536954 -10.393322 -8.395591 -8.770291

[[2]]
[1] -0.2173831  0.4013440 -2.2716535 -1.5199591  0.2817534

[[3]]
[1]  8.217122  7.209396  9.051645 10.628293 11.263494
```

Iterate over multiple objects **map2()** and **pmap()**:

```
1 mu <- c(-10, 0, 10)
2 sigma <- c(0, 0.1, 0)
3 purrr::map2(.x = mu, .y = sigma, .f = rnorm, n = 5)
```

```
[[1]]
[1] -10 -10 -10 -10 -10

[[2]]
[1] -0.097381610 -0.104544473 -0.063586494 -0.002777205
-0.105194011

[[3]]
[1] 10 10 10 10 10
```

For more details and variants see Advanced R [chapters 9-11](#) on functional programming.

# Easy parallelisation with furrr

- `{parallel}` and `{futures}` allow parallel coding over multiple cores.
- Powerful, but steep learning curve.
- `{furrr}` makes this very easy, just add `future_` to purrr verbs.

```
1 mu <- c(-10, 0, 10)
2 furrr::future_map(
3   .x = mu,
4   .f = rnorm,
5   .options = furrr::furrr_options(seed = TRUE),
6   n = 5)
```

```
[[1]]
[1] -11.210327 -7.806065 -10.330897 -9.628868 -11.735650

[[2]]
[1]  0.1394846 -2.5554590  0.8892910  1.0211497  1.4515885

[[3]]
[1] 10.22833  9.81417 10.17171 11.94347  9.74026
```

Need to be very careful handling RNG. See [R-bloggers](#) for more details.

# Sometimes R doesn't cut it



- An API for running C++ code in R
  - Loops that need to be run in order
  - Lots of function calls (e.g. deep recursion)
  - Fast data structures
- Beyond our scope but good to know exists. Starting point: Advanced R [Chapter 25](#).

# Wrapping up

## Summary

1. Pick you battles wisely
2. Target your energy with profiling
3. Scale loops with vectors
4. Scale loops in parallel processing
5. Scale in another language

## Help!

- Articles and blog links
- The R inferno ([Circles 2-4](#))
- Advanced R ([Chapters 23-25](#)),
- Efficient R ([Chapter 7](#)).

